

# Write-Once File System

George Bernard

Nick Lockett

Ryan St. Pierre

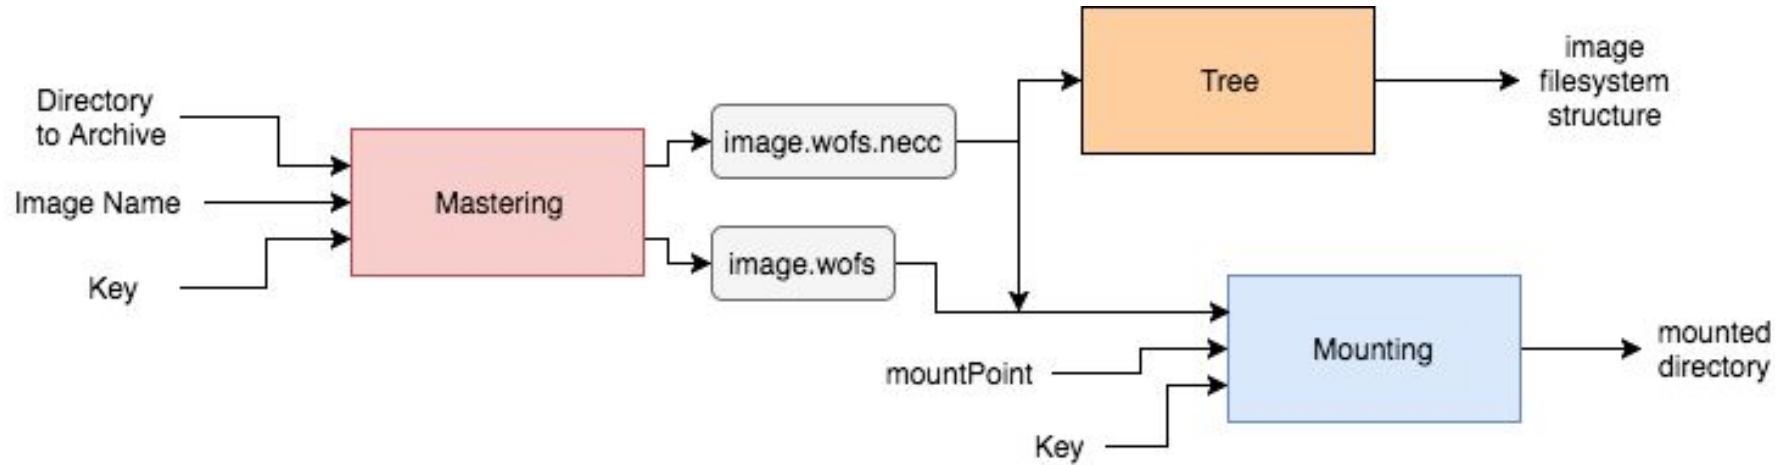
Matthew Wu

# Changes Since Project Presentation


- ❖ Removed key from command line arguments
- ❖ Allowed override when data integrity problems are detected
- ❖ Decoded status reporting
- ❖ Benchmarking and performance
- ❖ Code cleanup

# Technical Overview

# High Level Architecture




# WOFS Service Layered Architecture



'Onion' diagram

# On Disk Structure



Metadata header:

- ❖ Name
- ❖ Length
- ❖ Offset
- ❖ Time (of creation)
- ❖ Type (file or directory)

# Tradeoffs/Design Decisions

- ❖ Proprietary vs. known standard (e.g. .iso)
- ❖ Variable vs. fixed offset lists
- ❖ Metadata context
- ❖ File name size (space vs. flexibility)
- ❖ HMAC + ECC vs. keyed ECC (layered architecture)
- ❖ Programming language


# Mastering



Image the File System

Calculate and Append Hash

Apply ECC



# Mastering Code Structure

```
int run(std::string, std::string, std::string);

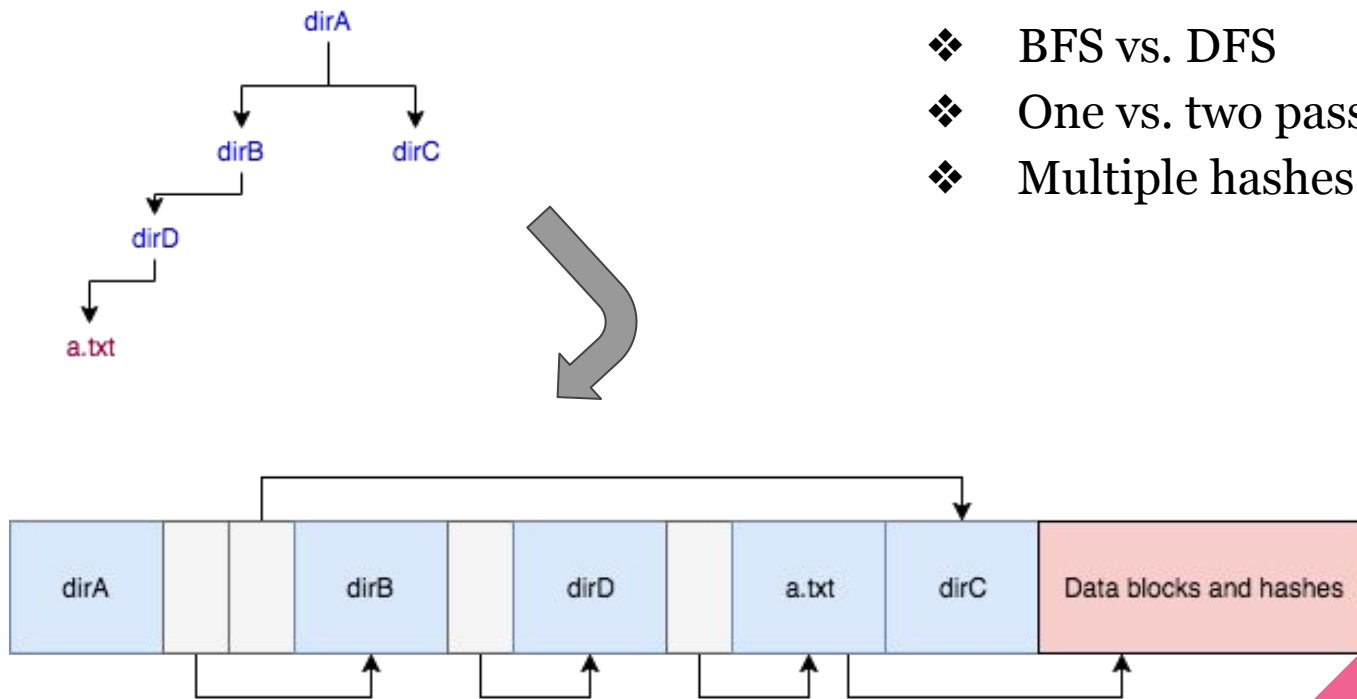
//Method applied on each node of the FS, required by nftw
static int s_builder(const char *, const struct stat *, int, struct FTW *);

//Major structural methods defining each stage of the process, in order of their usage
int imageDFS(const std::string& out_filename, node* root);
uint64_t writeDFS(node* node, FILE* output);
int hashAndAppend(const char*, const char*);
int addReedSolomon(std::string ifs, std::string ofs);

// Helper Methods
std::string parse_name(const std::string& path_name);
std::string space_pad(const std::string& s);
uint64_t find_header_size();
```

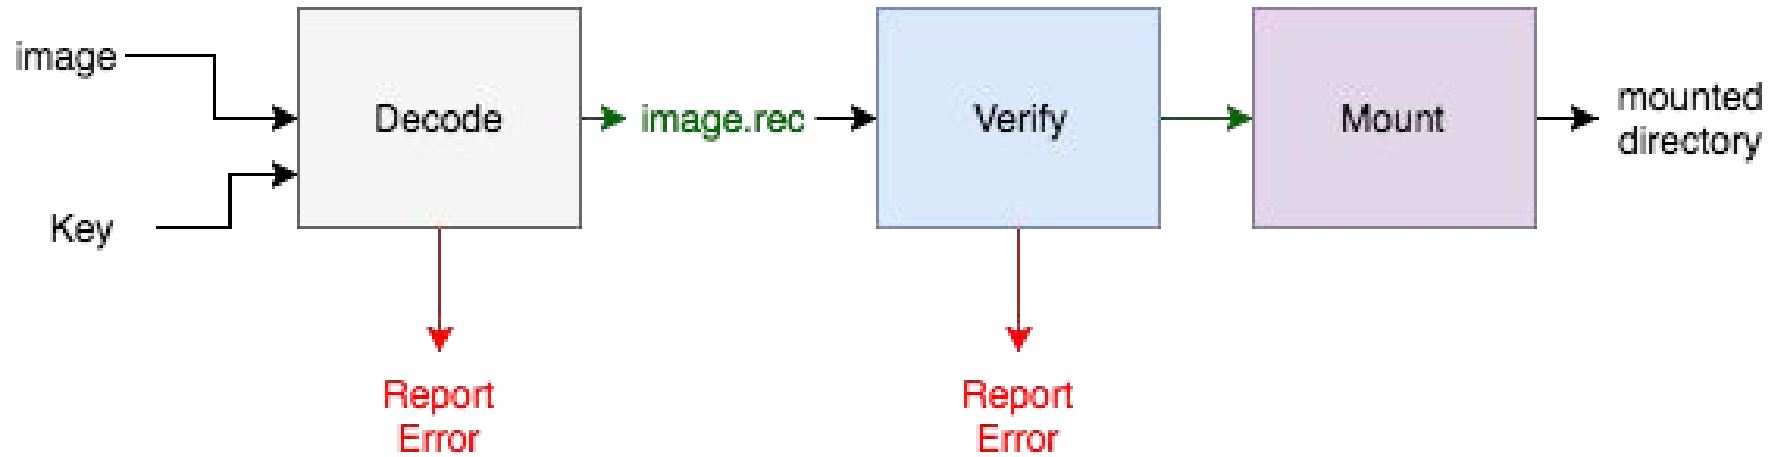
# Tradeoffs/Design Decisions

- ❖ BFS vs. DFS
- ❖ One vs. two passes
- ❖ Multiple hashes



# Link Blind

# Mounting



Decode the image (containing ECC)

Verify Hash

Mount image (FUSE)

# FUSE

- ❖ Chose FUSE commands for structure traversal, reading, and controlled access:
  - `getattr`
  - `readdir`
    - Current implementation is *path* based – should be offset based
  - `open`
  - `read`
- ❖ Fuse has ability to handle links (*symlink*, *unlink*, and *link*), but capability is not supported by mastering
- ❖ Leveraged FUSE's command line parser

# Mounting Code

```
// Return the metadata in the image file specified by path
static m_hdr* find(const char* path);

// Fill the attributes of stbuf for the file or directory and the given path
static int mount_getattr(const char *path, struct stat *stbuf,
                        struct fuse_file_info *fi);

// Place the contents of a directory into the buffer buf
static int mount_readdir(const char *path, void *buf, fuse_fill_dir_t filler,
                        off_t offset, struct fuse_file_info *fi,
                        enum fuse_readdir_flags flags);

// Verify that a file can be opened at the given path
static int mount_open(const char *path, struct fuse_file_info *fi);

//Return the content of the file at path into buf (at the given offset)
static int mount_read(const char *path, char *buf, size_t size, off_t offset,
                    struct fuse_file_info *fi);
```



# Mounting Limitations and Comments

## *Limitations*

### ❖ Size Constraints

- Fuse Block constraint (131,072 bytes reads)

- <https://github.com/libfuse/libfuse/issues/91>

- Reed-Solomon data word size (on decode)

- Fix: Caching unfinished read offset will speed traversal, but not limit # read calls

### ❖ .git and HEAD files

- Fix: Issue with git

# Mount + Master Demo

Tensorflow ML Library

<https://github.com/tensorflow/tensorflow>

- 993 Directories
- 9419 Files
- Image Size: 225MB



# ECC



# Motivations for ECC and HMAC

## **Archival Constraints:**

- ❖ Cold Time
- ❖ Copies from place to place
- ❖ Uses for Embedded Systems
- ❖ Lots of access beyond your organization

## **Fault Model:**

- ❖ Data corruption on a block level (not corrected by hardware)
- ❖ Any data corruption on an embedded system without ECC in place

## **Threat Model:**

- ❖ Malicious changes made between mastering and mounting

## **Reed Solomon:**

- ❖ Blocks of 256B, 32 Error Correcting bits
- ❖ Correct up to 16 byte errors per block
- ❖ Selected as a balance point between space and time efficiency

## **SHA-256 HMAC:**

- ❖ Private key used to verify file integrity
- ❖ Ensures no malicious editing, and that ECC returns correct data
- ❖ ‘Novice Proof’ access restriction

# ECC Block Size

# Reed Solomon

## Implementation:

- ❖ We used the Schifra library to implement Reed Solomon Codes
- ❖ (255, 223)
- ❖ Can correct up to 16 byte errors
- ❖ Optional

## ECC Tradeoffs:

- ❖ Flexibility
- ❖ Robustness
- ❖ Block Level

## RS Tradeoffs:

- ❖ Time Efficiency
- ❖ Space Efficiency
- ❖ Correction Bit

Overhead	Command	Shared Object	Symbol
14.60%	master.out	master.out	[.] schifra::galois::field::mul
11.73%	master.out	master.out	[.] schifra::galois::field_polynomial::operator[]
11.55%	master.out	master.out	[.] std::vector<schifra::galois::field_element, std::allocator<schifra::galois::field_ele
9.91%	master.out	master.out	[.] schifra::galois::field_polynomial::operator%=
9.58%	master.out	master.out	[.] schifra::galois::field_element::field_element
7.52%	master.out	master.out	[.] std::vector<schifra::galois::field_element, std::allocator<schifra::galois::field_ele
6.35%	master.out	master.out	[.] schifra::galois::field_element::operator=
4.37%	master.out	master.out	[.] schifra::galois::field_element::operator*=
3.65%	master.out	master.out	[.] schifra::galois::field_polynomial::operator[]
3.45%	master.out	master.out	[.] schifra::galois::operator+
3.17%	master.out	master.out	[.] schifra::galois::operator*
2.82%	master.out	master.out	[.] std::vector<schifra::galois::field_element, std::allocator<schifra::galois::field_ele
2.53%	master.out	master.out	[.] schifra::galois::field_element::operator+=
1.96%	master.out	master.out	[.] schifra::galois::field_element::~field_element
1.09%	master.out	master.out	[.] schifra::galois::field::div
0.37%	master.out	master.out	[.] std::construct<schifra::galois::field_element, schifra::galois::field_element const&
0.29%	master.out	master.out	[.] std::_Destroy<schifra::galois::field_element>
0.29%	master.out	master.out	[.] std::_Destroy_aux<false>::_destroy<schifra::galois::field_element*>
0.27%	master.out	master.out	[.] std::_addressof<schifra::galois::field_element>
0.26%	master.out	master.out	[.] schifra::galois::field_polynomial::deg
0.17%	master.out	master.out	[.] __gnu_cxx::operator!=<schifra::galois::field_element*, std::vector<schifra::galois::f
0.16%	master.out	[unknown]	[k] 0xffffffff8182d565
0.15%	master.out	master.out	[.] schifra::reed_solomon::file_encoder<255ul, 32ul, 223ul>::process_block
0.15%	master.out	master.out	[.] __gnu_cxx::__normal_iterator<schifra::galois::field_element*, std::vector<schifra::ga
0.14%	master.out	master.out	[.] __gnu_cxx::operator!=<schifra::galois::field_element const*, std::vector<schifra::gal
0.14%	master.out	master.out	[.] operator new
0.14%	master.out	master.out	[.] schifra::galois::field_element::operator=
0.13%	master.out	[unknown]	[k] 0xffffffff813fddd2
0.13%	master.out	master.out	[.] std::__uninitialized_fill_n<false>::_uninit_fill_n<schifra::galois::field_element*,
0.13%	master.out	master.out	[.] schifra::galois::field::size
0.12%	master.out	[unknown]	[k] 0xffffffff810a975d
0.11%	master.out	master.out	[.] std::forward<schifra::galois::field_element const&>
0.10%	master.out	master.out	[.] schifra::galois::operator/
0.10%	master.out	master.out	[.] schifra::galois::field_element::operator/=
0.09%	master.out	libc-2.23.so	[.] _int_malloc
0.08%	master.out	master.out	[.] schifra::reed_solomon::encoder<255ul, 32ul, 223ul>::msg_poly
0.08%	master.out	master.out	[.] std::__uninitialized_copy<false>::_uninit_copy<__gnu_cxx::__normal_iterator<schifra:
0.07%	master.out	master.out	[.] __gnu_cxx::__normal_iterator<schifra::galois::field_element const*, std::vector<schif
0.07%	master.out	master.out	[.] __gnu_cxx::__normal_iterator<schifra::galois::field_element const*, std::vector<schif
0.07%	master.out	master.out	[.] __gnu_cxx::__normal_iterator<schifra::galois::field_element*, std::vector<schifra::ga
0.06%	master.out	master.out	[.] std::_Destroy_aux<false>::_destroy<__gnu_cxx::__normal_iterator<schifra::galois::fie
0.06%	master.out	libc-2.23.so	[.] _int_free

For a higher level overview, try: perf report --sort comm,dso



# 95%

CPU Usage for Reed-Solomon ECC



# Reed Solomon Timing Estimates

- ❖  $O(n^2)$  runtime on the block size
- ❖  $O(n)$  runtime on the size of the file
- ❖ 4.0 GHz processor

An Example: 225MB file

- ❖  $C \cdot (223)^2 * (225\text{MB} \div 223\text{B}) \div (4\text{ GHz}) = 12.5 * C$  seconds
- ❖ This process took ~200 seconds

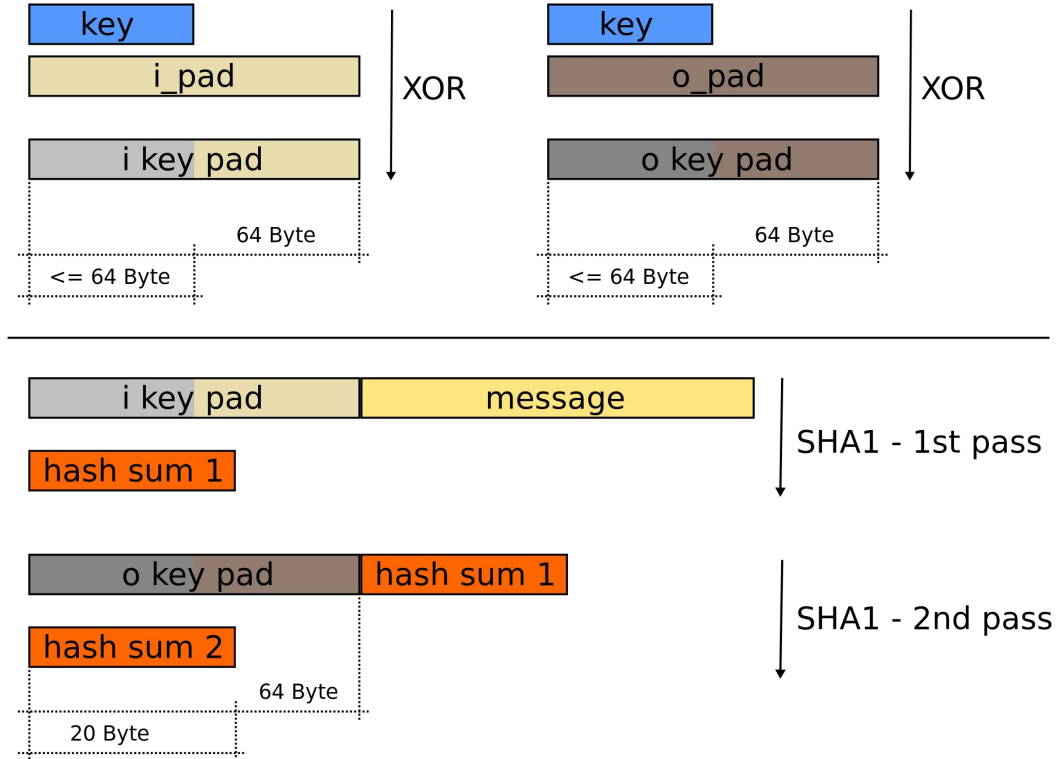
# Industry Solutions to Reed Solomon Slowness

- ❖ Hardware implementation
- ❖ Small Block size
- ❖ Less Error correcting bits

# Security

# SHA-256 HMAC

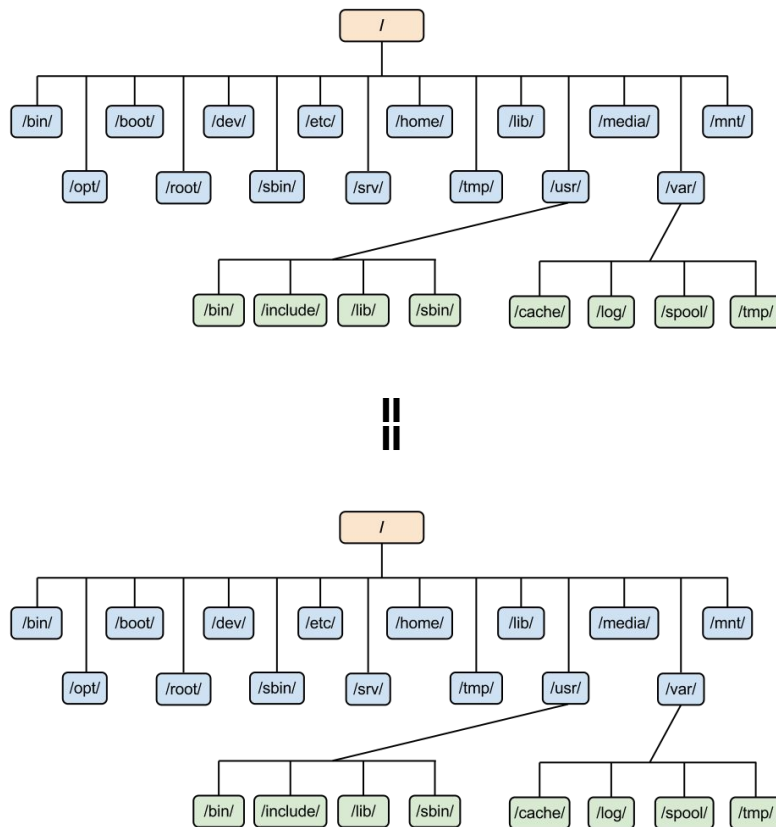
- ❖ OpenSSL library
- ❖ Hashed in blocks
- ❖ Shared private key between parties
- ❖ Salting is unnecessary given block sizes



# Testing

# ‘Mutation’ Testing

- ❖ Python script
- ❖ Comparison of metadata and file content
- ❖ Random shuffle access
- ❖ Specify how many trials



# Test Demo

`stress-test.py`

## Flags:

- ❖ `--trials`: Number of trials
- ❖ `--verbose`: Show each test output
- ❖ `--content`: Compare file content
- ❖ `--randomize`: Shuffle access to each file


# Benchmarking



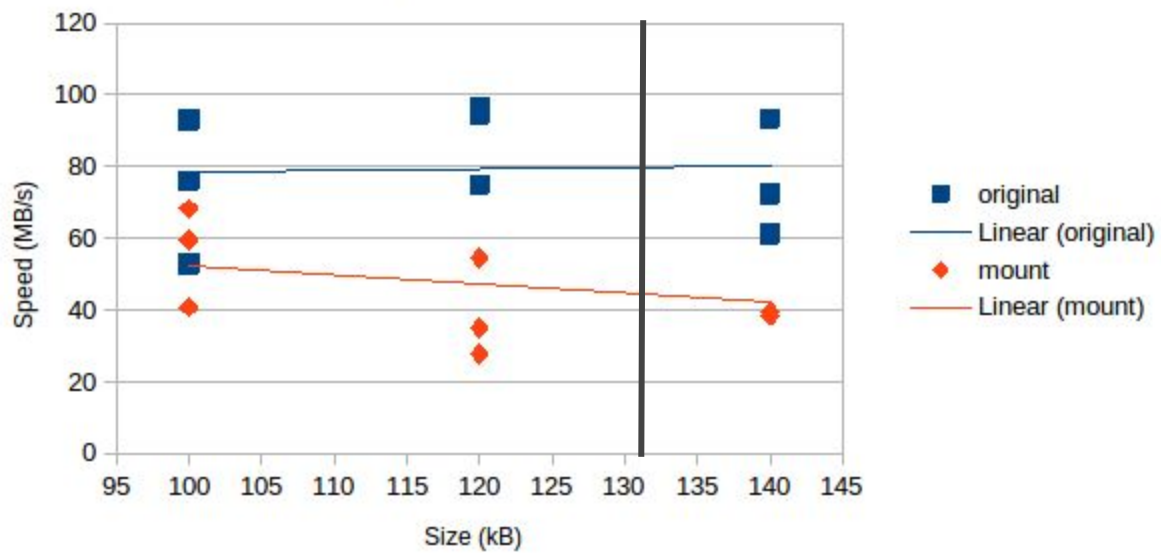
# Benchmarking

- ❖ Difficult to do in practice
  - Many benchmark programs create their own files (e.g. fio, bonnie++)
  - `fio` needed `unlink`
- ❖ Rely on linux staples
  - `time`
  - `dd`
- ❖ Profiling of stress test script
- ❖ SQLite also had difficulties


Sequential Read: Time Vs. Size



Sequential Read: Speed vs Time



Sequential: Speed vs. Size



# Interactive Q&A