

Write-Once File System

George Bernard

Nick Lockett

Ryan St. Pierre


Matthew Wu

Overview

Write-Once File System

Our project goal is to implement an archival, highly available write-once file system. Major changes since the proposal include:

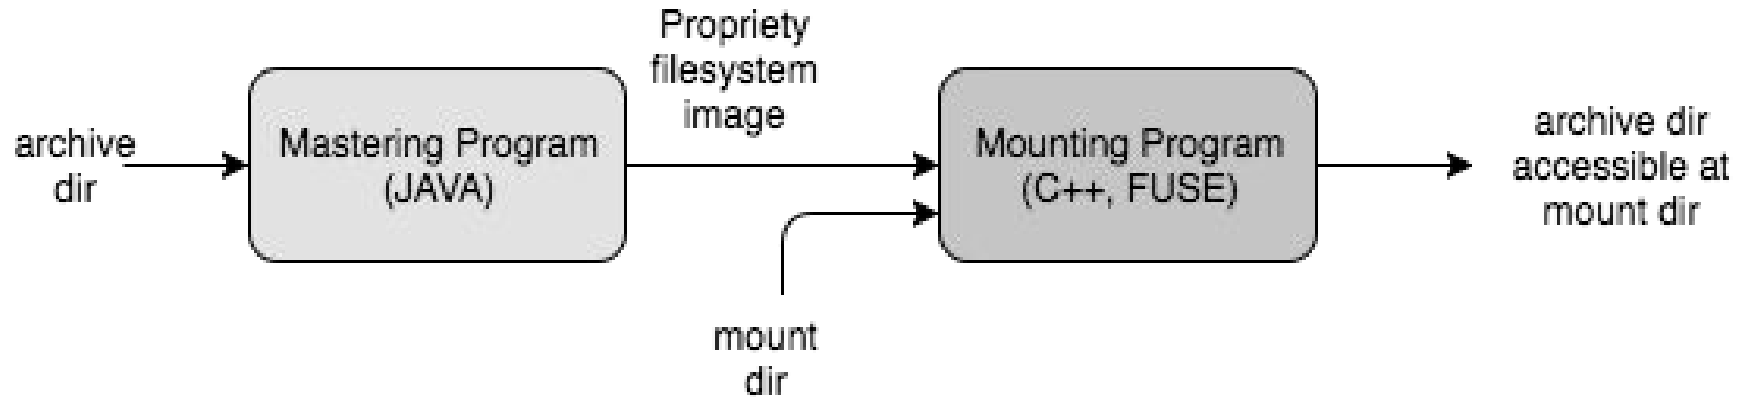
Major Changes

- ❖ Decision to use a proprietary image format for our filesystem rather than to meet .iso or other archive standards.
 - ❖ Changes to our extension goals:
 - Previously, our stretch goals were compression and bit level error correcting code.
 - During our proposal we found this does not make sense: space efficiency should not be an issue and bit level error correcting code should be handled by our program.
 - Instead, we will focus on block level failures, availability, information security, and ease of data extraction.
- 

Project Architecture

Write-Once Filesystem

General Architecture



Mastering Utility

- ❖ *Language:* Java
- ❖ *Platform:* Linux
- ❖ *Algorithm:*
 - Two Passes of filesystem *'tree'*
 - First, build graph based on existing filesystem
 - Second, traverse metadata graph, open files, and write them (along with the tree structure) to the image file.
- ❖ *File System File Representation:* Proprietary



Mounting Utility

- ❖ *Language:* C/C++
- ❖ *Platform:* Linux
- ❖ *Libraries:* FUSE (Filesystem in userspace)
- ❖ *Algorithm:* Pass in the mastered file, mount using FUSE, implement the read and directory methods.
 - Each call to FUSE traverses from the root of the mastered image file along the path given to the function



Timeline

Write-Once Filesystem

Milestone 1: Getting Started - October 6

- ❖ *Deliverables:*
 - Well defined structure for the filesystem image file.
 - Data structure to house the mastering information.
- ❖ *Group Members:* Since design is essential to the success of this project, all members will work together on this.
- ❖ *Server Usage:* Occasional light usage to examine file metadata and Unix protocols.



Milestone 2: Implement Mastering - October 15

- ❖ *Deliverable:*
 - Alpha implementation of WOFS mastering program
- ❖ *Group Members:* Leads – Nick and Matthew
- ❖ *New Skills:*
 - Java *File* class
- ❖ *Server Usage:* Light server usage for testing, most likely in the evening from 6 to midnight. Most development will be done in Java on local machine.



Milestone 3: Implement Mounting - October 22

- ❖ *Deliverable:*
 - Alpha implementation of WOFS mounting program
- ❖ *Group Members:* Leads – Ryan and George
- ❖ *New Skills:*
 - Libfuse, FUSE
- ❖ *Server Usage:* Significant usage for testing purposes, most likely in the evening from 6 to midnight.



Milestone 4: Milestone Presentation - October 25

- ❖ *Deliverable:*

- An integrated, demonstrable implementation of WOFS (mastering and mounting)
- 5 minute demo highlighting our current status

- ❖ *Group Members:* All



Milestone 5: Testing - Mid November

❖ *Deliverables:*

- Integration test pipeline for maintaining compatibility between both mastering and mounting programs, but also legacy versions of both programs
- Unit test suite for mastering and mounting programs.
- Beta version of mastering and mounting programs with bugs removed from alpha version.

❖ *Division of Labor:*

- Ryan and George: mounting unit and integration testing.
 - Matthew and Nick: mastering unit testing.
- 

Testing

- ❖ *Unit Testing frameworks:* JUnit (Java), Catch (C++)
- ❖ *Integration Tests:* We will need to test the interface between the master and mounting program.
 - This will likely be implemented as a *bash* script, or set of *bash* scripts, that will automatically input a test directory into the mastering program, mount the resulting file system image, and monitor the output of different linux commands on the mounted file system.
- ❖ *Performance:* We anticipate manually benchmarking our mastering and mounting utilities through the development process as user experience is impacted.

Milestone 6: Wrapping Up - November 31

- ❖ *Flexible time for any extensions or completion of any other deliverables*



Milestone 7: Final Presentation - December 7

❖ *Deliverables:*


- Finished, fully polished final product with accompanying test suit
- 15-minute rehearsed presentation



Criteria of Success

Write-Once Filesystem

Primary Criteria

- ❖ Functional correctness (i.e. it works). More formally, this means the file system is capable of mastering, mounting, and reading.
 - Master program creates a consistent and independent image file.
 - Master program capable of mastering directories with any file types and with tree structures including the following linux constructs: soft and hard links.
 - Mounting program is successfully able to mount any properly generated image file from the mastering program.
 - Correctly notify users of warning, errors, and failures of the program.
- 

Additional Criteria

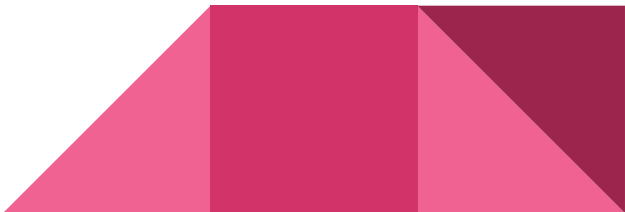
- ❖ Completing all milestones in a timely manner.
- ❖ Space and time efficiency of the file system.
 - No observable latency while using mounted file system.
- ❖ Full testing suite:
 - 80% code coverage.
 - Unit and integration tests.
- ❖ Extremely successful if we implement one stretch goal



Extensions

Write-Once Filesystem

Distributed File System with Fault Tolerant Reads

- ❖ A fitting extension for an archival system like this could be to distribute the image file across multiple machines with a corresponding checksum.
 - ❖ If an error is encountered, or if checksums differ across image files, the failed image can be reconstructed across the network.
 - ❖ In a complete implementation, if a failure is encountered while parsing the image file, the data could be requested across the network, ensuring availability with a performance hit.
 - ❖ The images could then have their checksums compared and the failed one could be copied from valid images.
- 

Encryption

- ❖ Following a Encrypt-Then-MAC procedure, the file system would be authenticated before it was unencrypted, keeping the underlying filesystem secure.
- ❖ Since any enterprise system should never roll their own encryption services, interfacing with a field tested and verified encryption library would be necessary.



Tree Command

This extension would be smaller than the other two and involves writing a program that traverses our file image and displays its structure before mounting. This would be a useful program for the user and also serve as a testing tool for the mastering program.



Major Risks and Challenges

- ❖ Learning Fuse/other tools: With any project, learning new tools essential to the project's success, such as Fuse, presents a risk when attempting to meet a strict deadline.
 - ❖ Dependency between mounting and mastering programs could present difficulties in the case that mastering takes longer than expected or does not work as expected, since mounting depends on a fully functioning mastering program.
 - ❖ Significant reliance on a single, early defined interface presents challenges if it turns out that the interface is actually not that well-suited for the project.
- 